

Hamiltonian and Sequential Monte Carlo An Ecosystem Example

Dominic Steinitz

27th September 2016

Outline

1 Typical Problems

Outline

- 1 Typical Problems
- 2 The Stochastic Model

Outline

- 1 Typical Problems
- 2 The Stochastic Model
- 3 LibBI

Outline

- 1 Typical Problems
- 2 The Stochastic Model
- 3 LibBI
- 4 Stan

Outline

- 1 Typical Problems
- 2 The Stochastic Model
- 3 LibBI
- 4 Stan
- 5 Thoughts

Some Application Areas

Suppose you have a model described by differential equations, e.g.

- Epidemiology: susceptible, infected, recovered (SIR)
- Pharmokinetics / pharmacodynamics (PK / PD)
- Ecology: predator / prey
- How to fit parameters?
- How to capture second order effects?

Dynamical Systems

Let X be a set (usually taken to be a complete metric space), T a monoid and $S : T \times X \rightarrow X$ a family of operators, often written as $S_t(x) \triangleq S(t, x)$.


A dynamical system is a triple of such objects (X, S, T) such that the semi-group property holds.

$$\begin{aligned} S_0(x) &= x \\ S_t(S_s(x)) &= S_{t+s}(x) \end{aligned}$$

S is called the evolutionary operator and sometimes the family $\{S_{t \in T}\}$ is called the flow of the dynamical system.

- Ordinary Differential Equation
- Delay Differential Equation
- Finite State Markov Chain
- Continuous State Markov Chain
- The Kalman Filter

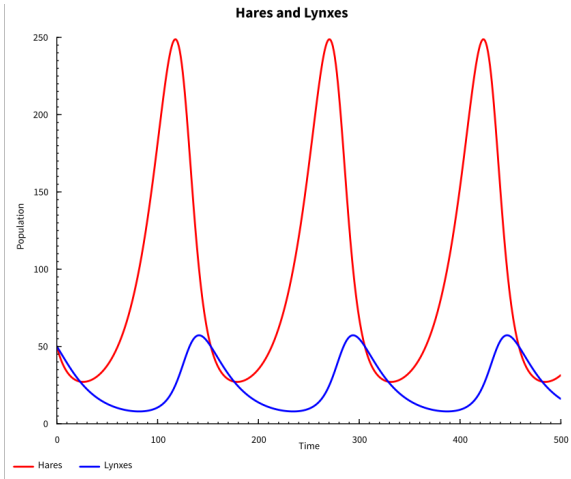
Hudson Bay Hares and Lynxes



Lotka and Volterra Original Model

$$\begin{aligned}\frac{dN_1}{dt} &= \rho_1 N_1 - c_1 N_1 N_2 \\ \frac{dN_2}{dt} &= c_2 N_1 N_2 - \rho_2 N_2\end{aligned}$$

A Typical Path

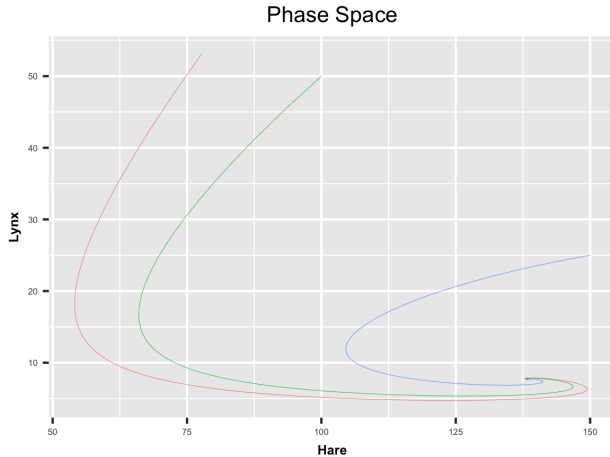


Structurally Unstable

- So it seems that this might be a good model.
- But can hare population really grow without a limit?
- Let us add carrying capacities.

$$\begin{aligned}\frac{dN_1}{dt} &= \rho_1 N_1 \left(1 - \frac{N_1}{K_1}\right) - c_1 N_1 N_2 \\ \frac{dN_2}{dt} &= -\rho_2 N_2 \left(1 + \frac{N_2}{K_2}\right) + c_2 N_1 N_2\end{aligned}$$

Typical Paths



Markov Transition Kernels

Recall that a **transition kernel** is a mapping $\mu : X \times \mathcal{Y} \rightarrow \overline{\mathbb{R}}_+$ where (X, \mathcal{X}) and (Y, \mathcal{Y}) are two measurable spaces such that $\mu(s, \cdot)$ is a probability measure on \mathcal{Y} for all $s \in X$ and such that $\mu(\cdot, A)$ is a measurable function on X for all $A \in \mathcal{Y}$.

Recall further that a family of such transitions $\{\mu_t\}_{t \in T}$ where T is some index set satisfying

$$\mu_{t+s}(x, A) = \int_Y \mu_s(x, dy) \mu_t(y, A) \quad (1)$$

gives rise to a Markov process.

Dynamical Systems as Markov Processes

A deterministic system can be formulated as a Markov process with a particularly simple transition kernel given by

$$\mu_t(x_s, A) = \delta(f_t(x_s), A) \triangleq \begin{cases} 1 & \text{if } f_t(x_s) \in A \\ 0 & \text{if } f_t(x_s) \notin A \end{cases} \quad (2)$$

where f_t is the deterministic state update function (the flow) and δ is the Dirac delta function.

Brownian Motion

- Now suppose deterministic system depends on some time-varying values.
- E.g. predator-prey model where the parameters cannot explain every aspect.

$$\mu_t(\theta_s, d\phi) = \mathcal{N}(d\phi | \theta_s, \tau^2(t-s)) \quad (3)$$

- $d\phi$ indicate probability densities.
- Parameter wiggles around like Brown's pollen particles rather than remaining absolutely fixed.

Ornstein-Uhlenbeck

- With Brownian motion, the parameters could drift off to $\pm\infty$.
- Parameters tend to stay close to some given value (mean-reverting).
-

$$\mu_t(\theta_s, d\phi) = \mathcal{N}\left(d\phi \mid \alpha + (\theta_s - \alpha)e^{-\beta t}, \frac{\sigma^2}{2\beta}(1 - e^{-2\beta t})\right)$$

- β expresses how strongly parameter responds to perturbations.
- α is the mean to which the process wants to revert (aka the asymptotic mean).
- σ^2 expresses how noisy the process is.

Stochastic Differential Equations

It is sometimes easier to view these transition kernels in terms of stochastic differential equations. Brownian motion can be expressed as

$$dX_t = \sigma dW_t$$

and Ornstein-Uhlenbeck can be expressed as

$$dX_t = -\beta(X_t - \alpha)dt + \sigma dW_t$$

where W_t is the Wiener process.

Capturing our Lack of Knowledge

Hare growth parameter

- Uncertainty grows over time
- Positive

$$\frac{dN_1}{dt} = \rho_1 N_1 \left(1 - \frac{N_1}{K_1} \right) - c_1 N_1 N_2$$

$$\frac{dN_2}{dt} = -\rho_2 N_2 \left(1 + \frac{N_2}{K_2} \right) + c_2 N_1 N_2$$

$$d\rho_1 = \rho_1 \sigma_{\rho_1} dW_t$$

W_t being a Wiener process.

Capturing our Lack of Knowledge II

By Itô we have

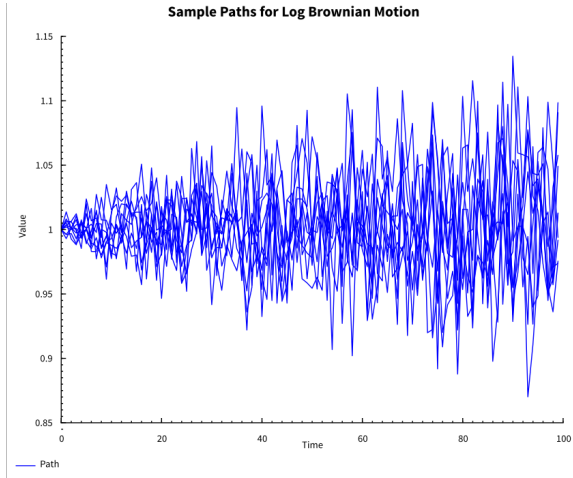
$$d(\log \rho_1) = -\frac{\sigma_{\rho_1}^2}{2}dt + \sigma_{\rho_1}dW_t$$

We can use this to generate paths for ρ_1 .


$$\rho_1(t + \Delta t) = \rho_1(t) \exp \left(-\frac{\sigma_{\rho_1}^2}{2} \Delta t + \sigma_{\rho_1} \sqrt{\Delta t} Z \right)$$

where $Z \sim \mathcal{N}(0, 1)$.

Log Brownian Paths



A Typical Noisy System Path



LibBI

- LibBI <http://libbi.org> is used for state-space modelling and Bayesian inference on high-performance computer hardware, including multi-core CPUs, many-core GPUs (graphics processing units) and distributed-memory clusters.
- Default method: PMMH / PMCMC C. Andrieu, A.D. & R. Holenstein, Particle Markov chain Monte Carlo for Efficient Numerical Simulation

LibBI Constants

```
model PP {  
  const h          = 0.1;      // time step  
  const delta_abs  = 1.0e-3;  // absolute error tolerance  
  const delta_rel  = 1.0e-6;  // relative error tolerance  
  
  const a  = 5.0e-1 // Hare growth rate - superfluous  
                // for inference but a reminder  
                // of what we should expect  
  const k1 = 2.0e2  // Hare carrying capacity  
  const b  = 2.0e-2 // Hare death rate per lynx  
  const d  = 4.0e-1 // Lynx death rate  
  const k2 = 2.0e1  // Lynx carrying capacity  
  const c  = 4.0e-3 // Lynx birth rate per hare
```

LibBI Parameters

```
state P, Z          // Hares and lynxes
state ln_alpha      // Hare growth rate (log of)
obs P_obs          // Observations of hares
param mu, sigma     // Mean and standard deviation of
                    // hare growth rate
noise w            // Noise
```



LibBI Initialisation

```
sub parameter {  
  mu ~ uniform(0.0, 1.0)  
  sigma ~ uniform(0.0, 0.5)  
}  
  
sub proposal_parameter {  
  mu ~ truncated_gaussian(mu, 0.02, 0.0, 1.0);  
  sigma ~ truncated_gaussian(sigma, 0.01, 0.0, 0.5);  
}  
  
sub initial {  
  P ~ log_normal(log(100.0), 0.2)  
  Z ~ log_normal(log(50.0), 0.1)  
  ln_alpha ~ gaussian(log(mu), sigma)  
}
```

LibBI Model

```
sub transition(delta = h) {  
  w ~ normal(0.0, sqrt(h));  
  ode(h = h, atoler = delta_abs, rtoler = delta_rel, alg =  
    'RK4(3)') {  
    dP/dt = exp(ln_alpha) * P * (1 - P / k1) - b * P * Z  
    dZ/dt = -d * Z * (1 + Z / k2) + c * P * Z  
    dln_alpha/dt = -sigma * sigma / 2 - sigma * w / h  
  }  
}  
  
sub observation {  
  P_obs ~ log_normal(log(P), 0.1)  
}  
}
```

Posteriors



Stan

Stan implements gradient-based Markov chain Monte Carlo (MCMC) algorithms for Bayesian inference, stochastic, gradient-based variational Bayesian methods for approximate Bayesian inference, and gradient-based optimization for penalized maximum likelihood estimation.

Stan implements reverse-mode automatic differentiation to calculate gradients of the model, which is required by HMC, NUTS, L-BFGS, BFGS, and variational inference. The automatic differentiation within Stan can be used outside of the probabilistic programming language.

Stan Constants

```
data {
  int<lower=1> T;      // Number of observations
  real y[T];          // Observed hares
  real k1;             // Hare carrying capacity
  real b;              // Hare death rate per lynx
  real d;              // Lynx death rate
  real k2;             // Lynx carrying capacity
  real c;              // Lynx birth rate per hare
  real deltaT;        // Time step
}
```

```
data=list(T = length(rdata_PP$P_obs$value),
          y = rdata_PP$P_obs$value,
          k1 = 2.0e2,
          b = 2.0e-2,
          d = 4.0e-1,
          k2 = 2.0e1,
          c = 4.0e-3,
          deltaT = rdata_PP$P_obs$time[2] -
                  rdata_PP$P_obs$time[1]
        ),
```

Stan Parameters



```
transformed parameters {
  real<lower=0> p[T]; real<lower=0> z[T]; real rho[T];
  p[1] = p0; z[1] = z0; rho[1] = rho0;
  for (t in 1:(T-1)){
    p[t+1] = p[t] +
              deltaT * f1 (exp(rho[t]), k1, b, p[t], z[t]);
    z[t+1] = z[t] + deltaT * f2 (d, k2, c, p[t], z[t]);
    rho[t+1] = rho[t] +
               sigma * w[t] - 0.5 * sigma * sigma * deltaT;
  }
}
```

```
functions {
  real f1 (real a, real k1, real b, real p, real z) {
    real q;
    q = a * p * (1 - p / k1) - b * p * z; return q;
  }
  real f2 (real d, real k2, real c, real p, real z) {
    real q;
    q = -d * z * (1 + z / k2) + c * p * z; return q;
  }
}
```


Stan Model

```
model {  
  mu      ~ uniform(0.0, 1.0);  
  sigma   ~ uniform(0.0, 0.5);  
  p0      ~ lognormal(log(100.0), 0.2);  
  z0      ~ lognormal(log(50.0), 0.1);  
  rho0    ~ normal(log(mu), sigma);  
  w       ~ normal(0.0, sqrt(deltaT));  
  
  for (t in 1:T) {  
    y[t] ~ lognormal(log(p[t]), 0.1);  
  }  
}
```

Posteriors



Some Thoughts

- Designing programming languages especially probabilistic ones is hard.
- Stan and LibBI capture 30 years of advance in modelling but ignore 30 years of advance in programming languages.
- Both take time to compile and report errors.
- When things don't work, diagnosis is hard.
- Venture (née Church), Figaro, BLOG, JAGS, monad-bayes, PyMC(3)? and many more.